

Hidden Markov Models

Theory, Algorithms, and Applications

Alexey Kravatskiy

May 29, 2025

- 1 Introduction to Hidden Markov Models
- 2 Baum-Welch Algorithm
- 3 Viterbi Algorithm

What is a Hidden Markov Model?

- A statistical model that represents a system with hidden states
- Based on Markov processes where future states depend only on the current state
- Consists of:
 - Hidden states X_t (not directly observable)
 - Observable outputs Y_t
 - Transition probabilities between states
 - Emission probabilities for observations

Mathematical Definition

A discrete HMM is defined by:

- N hidden states: $S = \{1, 2, \dots, N\}$
- M possible observations: $V = \{v_1, v_2, \dots, v_M\}$
- Transition matrix $A = \{a_{ij}\}$ where $a_{ij} = P(X_{t+1} = j | X_t = i)$
- Emission matrix $B = \{b_j(k)\}$ where $b_j(k) = P(Y_t = v_k | X_t = j)$
- Initial state distribution $\pi = \{\pi_i\}$ where $\pi_i = P(X_1 = i)$

Baum-Welch Algorithm for hidden structure

- Also known as the Forward-Backward algorithm
- Used to find the maximum likelihood estimate of HMM parameters
- Given observation sequence $Y = (Y_1 = y_1, Y_2 = y_2, \dots, Y_T = y_T)$
- Finds $\theta^* = \arg \max_{\theta} P(Y|\theta)$ (or a stationary point)
- Parameters $\theta = (A, B, \pi)$ where:
 - $A = \{a_{ij}\}$: transition probabilities
 - $B = \{b_j(y_i)\}$: emission probabilities
 - π : initial state distribution

EM Principle

- Iterative method for finding maximum likelihood estimates
- Handles incomplete data by treating hidden states as missing data
- Each iteration consists of two steps:
 - E-step: Compute expected value of log-likelihood
 - M-step: Maximize this expectation

Forward-Backward Procedure

Forward Variable $\alpha_i(t)$

Probability of observing sequence y_1, \dots, y_t and being in state i at time t :

$$\alpha_i(t) = P(Y_1 = y_1, \dots, Y_t = y_t, X_t = i | \theta)$$

- Initialization:

$$\alpha_i(1) = \pi_i b_i(y_1)$$

- Recursion:

$$\alpha_j(t) = b_j(y_t) \sum_{i=1}^N \alpha_i(t-1) a_{ij}$$

- Termination:

$$P(Y|\theta) = \sum_{i=1}^N \alpha_i(T)$$

Forward-Backward Procedure

Backward Variable $\beta_i(t)$

Probability of observing sequence y_{t+1}, \dots, y_T given state i at time t :

$$\beta_i(t) = P(Y_{t+1} = y_{t+1}, \dots, Y_T = y_T | X_t = i, \theta)$$

- Initialization:

$$\beta_i(T) = 1$$

- Recursion:

$$\beta_i(t) = \sum_{j=1}^N a_{ij} b_j(y_{t+1}) \beta_j(t+1)$$

E-step: Computing Intermediate Variables

- Using Bayes' theorem and forward-backward variables:
 - Probability of being in state i at time t :

$$\gamma_i(t) = P(X_t = i | Y, \theta) = \frac{\alpha_i(t)\beta_i(t)}{P(Y|\theta)}$$

- Probability of transition from i to j :

$$\xi_{ij}(t) = P(X_t = i, X_{t+1} = j | Y, \theta) = \frac{\alpha_i(t)a_{ij}b_j(y_{t+1})\beta_j(t+1)}{P(Y|\theta)}$$

- These probabilities are used to compute expected counts:
 - Expected time spent in state i : $\sum_{t=1}^T \gamma_i(t)$
 - Expected transitions from i to j : $\sum_{t=1}^{T-1} \xi_{ij}(t)$

M-step: Parameter Updates

- Update parameters to maximize expected log-likelihood:

$$\begin{aligned}\bar{a}_{ij} &= \frac{\sum_{t=1}^{T-1} \xi_{ij}(t)}{\sum_{t=1}^{T-1} \gamma_i(t)} \\ \bar{b}_j(k) &= \frac{\sum_{t=1}^T \gamma_j(t) \cdot \delta(y_t, v_k)}{\sum_{t=1}^T \gamma_j(t)} \\ \bar{\pi}_i &= \gamma_i(1)\end{aligned}$$

- Each update increases the likelihood:

$$P(Y|\theta_{new}) \geq P(Y|\theta_{old})$$

Convergence Properties

- Baum-Welch is guaranteed to converge
- However, it may converge to:
 - Local maximum (most common)
 - Saddle point (rare)
 - Global maximum (not guaranteed)
- Quality of solution depends on:
 - Initial parameter values
 - Model structure
 - Amount of training data

Viterbi Algorithm for the hidden sequence

- Dynamic programming algorithm
- Finds most likely sequence of hidden states (Viterbi path)
- Given observation sequence $Y = (Y_1 = y_1, Y_2 = y_2, \dots, Y_T = y_T)$
- Uses two matrices of size $T \times N$:
 - $P_{t,s}$: maximum probability of ending at state s at time t
 - $Q_{t,s}$: previous state in the maximum probability path

Viterbi Algorithm Steps

- 1 Initialization ($t = 0$):

$$P_{0,s} = \pi_s \cdot b_s(y_0)$$

$$Q_{0,s} = 0$$

- 2 Recursion ($t > 0$):

$$P_{t,s} = \max_{r \in S} (P_{t-1,r} \cdot a_{r,s} \cdot b_s(y_t))$$

$$Q_{t,s} = \arg \max_{r \in S} (P_{t-1,r} \cdot a_{r,s})$$

- 3 Termination:

$$P^* = \max_{s \in S} P_{T-1,s}$$

$$X_{T-1}^* = \arg \max_{s \in S} P_{T-1,s}$$

- 4 Path backtracking:

$$X_t^* = Q_{t+1}(X_{t+1}^*)$$

Applications

- Speech recognition
- Natural language processing
- Bioinformatics
- Financial time series analysis
- Weather forecasting

Thank You!

Questions?